

Operating Systems & Concurrency: Process Concepts

Michael Brockway

October 6, 2011

Outline

- ▶ Processes - context, data area, states
- ▶ Process creation, termination – unix examples
- ▶ Processes and threads

Processes

Definition: A *process* is an *instance* of a program (an application) that is *running* under the management of the operating system.

A process runs *sequentially*

Modern operating systems *multitask* – they manage the simultaneous running of several (lots!) of application processes. The management is performed by a *scheduler*.

Processes

Each process has

- ▶ a process *ID*
- ▶ a *priority*
- ▶ its own *context* and *state* (see below)
- ▶ CPU *scheduling* information
- ▶ file handles, network ports; I/O status information
- ▶ a data area; memory management information
- ▶ etc

These data form the fields of the *process control block* (PCB), aka *task control block* (TCB) for the process/task. The operating system keeps a table of PCBs for all of the currently executing processes.

Process context

This is the set of values in all the *CPU registers* including

- ▶ the *program counter*: the address in memory from where the next instruction will be fetched
- ▶ the *program status register*: bits are set according to the result of the last operation; the current operation may test these bits
- ▶ the *stack pointer*: the address of the current top of the stack
- ▶ one or more data registers (*accumulators*) for holding data temporarily during a computation

Subroutine calls are normally managed by storing return address, parameters, return value(s) on the *stack*

Process data area

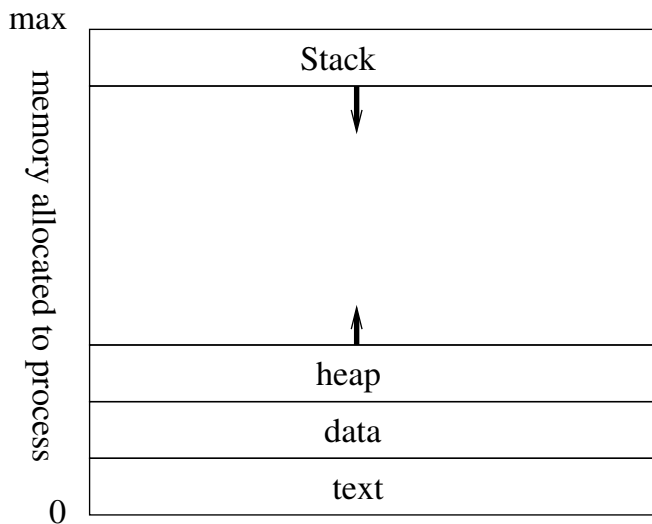


Figure: Data area for a process

Process states




Figure: Life cycle of states for a process

In a *multitasking* system, there will be several (lots) of tasks (processes) “running”. One will be actually *running* on the CPU; the others will be *ready*. The *scheduler* periodically performs a *context switch* to give all the tasks a turn. Usually this is rapid, giving the appearance of simultaneously, *concurrently* running processes.

The processes/tasks are *logically* concurrent.

If there are multiple CPUs, there will be multiple *running* tasks – 1 per CPU. In this case we have some actual *physical* concurrency.

Process creation

A “parent” process create “children” processes, which, in turn create other processes, forming a tree of processes.

Generally, a process is identified and managed via a *process identifier* (pid)

Resource sharing

- ▶ Parent and children share all resources; or
- ▶ Children share a subset of the parents resources; or
- ▶ Parent and child share no resources

Execution

- ▶ Parent and children execute concurrently; or
- ▶ Parent waits until children terminate

Process creation - UNIX example

The *fork* system call creates a new process

The *exec* system call used after a fork to replace the process memory space with a new program




Figure: Fork example

Process creation - UNIX example

```
int main() {
pid_t  pid;
    pid = fork();                /* fork another process */
    if (pid < 0) {                /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait (NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```

Processes termination (UNIX)

Process executes last statement and asks the operating system to delete it (exit)

- ▶ Output data from child to parent (via wait)
- ▶ Process resources are deallocated by operating system

Parent may terminate execution of child processes (abort)

- ▶ if child has exceeded allocated resources, or
- ▶ task assigned to child is no longer required, or
- ▶ if parent is exiting

Some operating systems do not allow child to continue if its parent terminates: all children are terminated - *cascading termination*.

More on fork

Notice from the code that `fork` returns an integer value -

- ▶ 0 if it is in the parent process
- ▶ a nonzero value in the child (clone) – in fact, the process ID of the parent.

If `exec()` is not used, `fork` simply creates a “clone” of the parent process: global variables, file handles etc are duplicated. Can you explain the behaviour of the following program (`processdemo.c` in Source files set 1, downloadable from [here](#))?

```
#include <stdio.h>
int x = 50;    /* a global variable */
```

More on fork

```
void adjustX(char * legend, int i)
{ long p;
  while (1) /* loop forever */
  { printf("%s: %i\n", legend, x);
    x += i;
    p=0;
    while (p<1000000000) p++; /* a "busy" delay */
  }
}

main()
{ int c;
  printf("creating new process:\n");
  c = fork();
  printf("process %i created\n", c);
  if (c==0)
    adjustX("child", 1); /* child process */
  else
    adjustX("parent", -1); /* parent process */
}
```

Processes and Threads

A “simple” process executes *sequentially* – one operation at a time: a *single-threaded process*

A process *may* be *multithreaded*

- ▶ Several threads, *each one* executes sequentially
- ▶ Each thread is scheduled as if it were a separate process
 - ▶ Each thread has its own subroutine stack
 - ▶ Each thread has a distinct state, its own scheduling information
- ▶ But the threads belonging to a particular process share all other data, memory management information, file handles, network ports, I/O status information

A thread is a “light-weight” process-like entity; *part of* a process.

Normally a process is ended when all its threads have ended.

- ▶ A *daemon* is a thread which carries on executing after its parent process has finished.

UNIX (POSIX) Threads example: `threaddemo.c`

The source file is in "Source files set 1", downloadable from [here](#).

```
#include <stdio.h>
#include <pthread.h>

int x = 50;    /* a global (shared) variable */

void * adjustX(void *n)
{  int i = (int)n;
   long p;
   while (1)    /* loop forever */
   {   printf("adjustment = %i; x = %i\n", i, x);
       x += i;
       p=0;
       while (p<1000000000) p++;    /* a "busy" delay */
   }
   return(n);
}
```

UNIX (POSIX) Threads example

```
main()
{  int a;
   pthread_t  up_thread, dn_thread;

   pthread_attr_t *attr; /* thread attribute variable */
   attr=0;

   printf("creating threads:\n");
   pthread_create(&up_thread,attr, adjustX, (void *)1);
   pthread_create(&dn_thread,attr, adjustX, (void *)-1);

   while (1) /* loop forever */
   { ;}
}
```


Operating Systems & Concurrency: Process Scheduling and Communication

Michael Brockway

October 6, 2011

Scheduling

The mix of running processes is managed by a *scheduler* process which gives the “running” processes turn-about on the CPU.

- ▶ This may be a straight “round-robin”; or
- ▶ it may be that processes are assigned *priorities*:
 - ▶ A higher priority process is given the CPU ahead of a lower-priority process.

The scheduler manages a set of processes/tasks that are *ready to run*. One of them is currently *running*. In a *context-switch* performed by the scheduler the running task changes places with one of the ready tasks.

This happens frequently (commonly 20-50 times per second) so that the processes/tasks appear to be running simultaneously.

Scheduling




Figure: States of a process

In reality a process passes between a number of *states*, doing many *context switches* before finishing.

When the process requires an operating system *resource* (graphics, disk or network I/O, or a communication from another process) in order to complete an operation it may have to *wait*.

The scheduler moves it from the *running* state to the *waiting* state. Eventually a waiting task “wakes up” – when it has the resource, communication or whatever it was waiting for – and return to the *ready* state: it is eligible again for context switch into *running* state.

Scheduling - process context

Each process has its own *context*:

- ▶ program counter;
- ▶ program status register;
- ▶ other hardware *registers*

Recall the *fetch-execute cycle* -

1. using address in program counter, fetch next instruction from memory;
2. increment program counter by size of instruction;
3. decode the instruction
4. execute the instruction
5. check program status register and possibly reload program counter
6. check for an *interrupt*

Scheduling – interrupts

The CPU can be *interrupted* by an event in its environment – a signal from a peripheral. It handles the interrupt by

1. saving the process context;
2. looking up the address of the interrupt's *handler* routine in the *interrupt vector table*;
3. running the handler
4. restoring the process context - this resumes the process

In a multitasking environment there is a *timer* which periodically fires an interrupt. The handler determines using a *scheduling algorithm* whether the currently running task has had a long enough turn on the CPU and if so, effect a context-switch: restore the context of *another* ready task. The former running task goes to the back of the ready queue: the dotted arrow in figure 1) is thus effected two transitions via the *interrupted* state.

In general the processes/tasks in the ready and waiting states are in *prioritized queues*.

Scheduling - flow of states

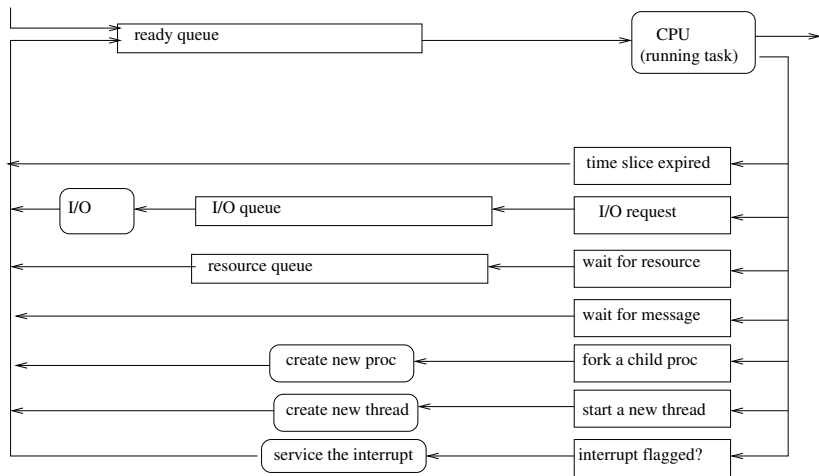


Figure: "Life cycle" flow

Scheduling activities

Long-term scheduling

- ▶ repeat period = seconds or minutes
- ▶ admits new processes to ready queue

Short-term scheduling

- ▶ repeat period = milliseconds
- ▶ triggers a context switch

I/O can take a long time – longer than the short-term scheduling interval. In this case the process is said to be *I/O-bound*.

At the opposite extreme is a process which spends most of its time performing computations - a *CPU-bound* process.

Scheduling Algorithms

How does the scheduler decide which process to run? Aim to -

- ▶ be fair – give all processes a fair amount of CPU time
- ▶ minimize response time
- ▶ minimize turn around time
- ▶ maximize CPU utilization
- ▶ meet user deadlines
- ▶ maximize system utilization

These can't ALL be satisfied!

Types of Scheduling Algorithms

A *non-preemptive* scheduler allows the *running* task to continue running until it gives up the CPU: changes state because it is *waiting* for a resource or message. Examples -

- ▶ first-come-first-served
- ▶ shortest-job-first

With a *preemptive* scheduler, the tasks/processes all have priorities assigned and the *running* task is switched out to *ready* as soon as a higher-priority task becomes *ready*. If its time-slice is exhausted before this happens another equal-priority task may be switch-in.

- ▶ Priorities are be assigned statically or dynamically
- ▶ “Round-robin” scheduling fits within this approach

Scheduling Examples

Scenario: 3 jobs:

1. loop (3ms CPU, 3 ms I/O time) 6 times
2. loop (1ms CPU, 5 ms I/O time) 6 times
3. loop (8ms CPU, 1 ms I/O time) 4 times




Figure: First-come first-served

(shaded = CPU time, unshaded = I/O pending)

Scheduling Examples - ctd




Figure: Shortest job first




Figure: Pre-emption, I/O-bound jobs have priority

Scheduling Examples - ctd




Figure: Pre-emption, CPU-bound jobs have priority

Scheduling - priorities

Priorities can be assigned to jobs *statically*

- ▶ In MicroC, a *task* is assigned a priority when it is created.


A option is to be able to change a task priority even when it is running.

Algorithms exist for assigning priorities for example -

- ▶ A *periodic* task has a “duty cycle”, a loop which repeats at a fixed interval. *Rate-monotonic scheduling* assigns priorities in *increasing* order of frequency – more frequently recurring tasks are given higher priority.
- ▶ A set of tasks have have *deadlines* defined. *Deadline-monotonic scheduling* assigns priorities in order of deadlines, the earlier the deadline, the higher the priority.
- ▶ A priority may be adjusted according to the “age” the

Priority Inversion

An example of what can go wrong! Consider three tasks (low, medium, high priority) and a shared resource, a data file being written to.



0. Initially the low priority task, 3, is running by itself ...
1. After a while it obtains an exclusive lock on a data file ...
2. Then high-priority task 1 becomes ready and pre-empts task 3 ...
3. But then it won't the file locked by 3, so must wait. 3 resumes ...
4. Task 2 (medium pri) pre-empts task 3 which still has file locked ...
5. Task 2 finishes, task 3 resumes. Task one still waiting for file.
6. Task 3 releases lock on file, so task 1 can obtain it and resume.
Task 1, *high priority* has been kept waiting a long time by the lower priority tasks!

Communication between Tasks

Processes within a system may cooperate: communicate or synchronise, affect one another, share data. Reasons include

- ▶ Information sharing
- ▶ Computation speedup
- ▶ Modularity
- ▶ Convenience

Interprocess communication (IPC) may be by *shared memory* or by *message passing*

Communication by Shared Memory

Shared memory may use a *bounded buffer*

- ▶ a shared variable, or
- ▶ a “circular” array
 - ▶ Indexes are incremented and wrap around when they reach the end

In case such as these *synchronisation* is required to prevent a writer process overwriting the buffer before a reader process has read data there, and to prevent a reader reading data already read, before it is refreshed.

- ▶ The producer-consumer problem – see `prodconsUnsync.c` in Source files set 1, downloadable from [here](#)

Communication by Message Passing

May be

- ▶ Synchronous (*blocking*) – sending (receiving) process *waits* until receiving (sending) process has “synchronised” or “rendezvoused”;
- ▶ Asynchronous (*non-blocking*) – eg
 - ▶ sending process puts message in a *message queue* or *mail box*
 - ▶ receiving process takes message from queue or mail box
 - ▶ sender may *wait* if queue/mail box full;
 - ▶ receiver may *wait* if queue/mail box empty.
- ▶ Unicast, multicast, broadcast
- ▶ unidirectional, bidirectional

You will apply a number of these techniques for synchronisation (including Dijkstra's *semaphores* and *mutexes*) and message passing in real-time embedded systems.