

Unix C API for TCP

kv5002

Dr Alun Moon

Computer Science

March 14, 2019

A note on man pages

How they are referred to

- Entries in the man pages are given with the section number in parenthesis.

quoted as	man command	
ip(7)	man 7 ip	man ip.7
udp(7)	man 7 udp	man udp.7
tcp(7)	man 7 tcp	man tcp.7
service(5)	man 5 service	man service.5

The UNIX API for the TCP/IP stack is well documented in the manual pages.

Manual sections

- 1 Executable programs or shell commands
- 2 System calls (functions provided by the kernel)
- 3 Library calls (functions within program libraries)
- 4 Special files (usually found in `/dev`)
- 5 File formats and conventions eg `/etc/passwd`
- 6 Games
- 7 Miscellaneous (including macro packages and conventions),
e.g. `man(7)`, `groff(7)`
- 8 System administration commands (usually only for root)
- 9 Kernel routines [Non standard]

Network byte order

`byteorder(3)`, `htons(3)`, `ntohs(3)`

- addresses are 32bit (in ip4)
- ports are 16bit
- *byte order matters*

Utility functions in `<arpa/inet.h>`,
standard types in `<stdint.h>`

- convert to and from network and host byte ordering
- work on 16bit (short) and 32bit (long) values

Address lookup

`getaddrinfo(3)`

```
#include <sys/types.h>
#include <sys/socket.h>
#include <netdb.h>

int getaddrinfo(const char *node, const char *service,
                const struct addrinfo *hints,
                struct addrinfo **res);

void freeaddrinfo(struct addrinfo *res);

const char * gai_strerror(int errcode);
```

getaddrinfo

- Given node and service, which identify an Internet host and a service, `getaddrinfo()` returns one or more `addrinfo` structures, each of which contains an Internet address that can be specified in a call to `bind(2)` or `connect(2)`.
- `getaddrinfo()` returns 0 if it succeeds, or a nonzero error code: The `gai_strerror()` function translates these error codes to a human readable string, suitable for error reporting.
- The `getaddrinfo()` function allocates and initializes a linked list of `addrinfo` structures, one for each network address that matches node and service, subject to any restrictions imposed by hints, and returns a pointer to the start of the list in `res`. The items in the linked list are linked by the `ai_next` field. The sorting function used within `getaddrinfo()` is defined in RFC 3484;

addrinfo structure

```
struct addrinfo {  
    int ai_flags;  
    int ai_family;  
    int ai_socktype;  
    int ai_protocol;  
    socklen_t ai_addrlen;  
    struct sockaddr *ai_addr;  
    char *ai_canonname;  
    struct addrinfo *ai_next;  
};
```

Field values

to set in hints

`ai_family` set to

`AF_INET` for ip v4 (see `ip(7)`)

`AF_INET6` for ip v6 (see `ipv6(7)`)

`AF_UNSPEC` for either

`ai_socktype` to

`SOCK_STREAM` for TCP sockets (see `tcp(7)`)

`SOCK_DGRAM` for UDP sockets (see `udp(7)`)

`ai_flags` combine flags by OR

`AI_PASSIVE` and `node` is `NULL`, then the returned socket addresses will be suitable for `bind(2)`ing a socket that will accept(2) connections. *Used in creating servers.*

`AI_CANONNAME` the `ai_canonname` field of the first of the `addrinfo` structures in the returned list is set to point to the official name of the host.

Example

set-up and query

```
struct addrinfo  hints;
struct addrinfo *results;
int err;

hints.ai_family = AF_UNSPEC;
hints.ai_socktype = 0;
hints.ai_protocol = 0;
hints.ai_flags = AI_CANONNAME;

err = getaddrinfo( "hesabu.net", "http", &hints, &results);

if( err) {
    fprintf(stderr, "Error trying to open %s:%s\n    %s\n",
              argv[1], argv[2], gai_strerror(err));
    exit(EXIT_FAILURE);
}
```

Example

handle results

```
while( results ) {
    struct sockaddr_in *ipaddr;
    ipaddr = (struct sockaddr_in *)results->ai_addr;
    printf("    canonical name: %s\n", results->ai_canonname);
    printf(" address : %s\n", inet_ntoa( ipaddr->sin_addr ) );
    printf("    port      : %d\n", ntohs(ipaddr->sin_port) );

    results = results->ai_next;
}
```

typically just use

```
struct sockaddr_in *ipaddr =
    (struct sockaddr_in *)results->ai_addr;
```

Sockets

socket(7), socket(2)

```
#include <sys/socket.h>
```

```
sockfd = socket(int socket_family, int socket_type,  
                int protocol);
```

socket() creates an endpoint for communication and returns a file descriptor that refers to that endpoint. The file descriptor returned by a successful call will be the lowest-numbered file descriptor not currently open for the process.

typical use

```
int sfd = socket(AF_INET, SOCK_STREAM, 0);  
if( sfd == -1 ) { /* error handling */ }
```

Connecting to a server

connect(2)

```
#include <sys/types.h>
#include <sys/socket.h>
int connect(int sockfd,
            const struct sockaddr *addr, socklen_t addrlen);
```

The `connect()` system call connects the socket referred to by the file descriptor `sockfd` to the address specified by `addr`. The `addrlen` argument specifies the size of `addr`. The format of the address in `addr` is determined by the address space of the socket `sockfd`; see `socket(2)` for further details.

Using results from getaddrinfo

The `getaddrinfo()` function has been used to look a server and port. The results are pointed to by `results`.

```
err = getaddrinfo( "hesabu.net", "http", &hints, &results);  
if( err == -1 ) { /* error handling */ }  
  
err = connect( sfd, results->ai_addr, results->ai_addrlen);  
if( err == -1 ) { /* error handling */ }
```

Creating a client using TCP

- Lookup server with name/ip address and port/service
`getaddrinfo(3)`
- create a socket `socket(2)`
- connect to the server `connect(2)`
- use `send(2)` and `recv(2)` to send and receive messages.

Example

```
int s = socket(AF_INET, SOCK_STREAM, 0);
if( s == -1){
    perror("Error creating socket");
    exit(EXIT_FAILURE);
}

int c = connect(s, results->ai_addr, results->ai_addrlen);
if( s == -1){
    perror("Error connecting to server");
    exit(EXIT_FAILURE);
}
```

Sending Messages

- declare buffer as `char buffer[BUFSIZE]`
- also add `char *message = buffer;`
- use formatted string writing
`sprintf(buffer, "formats..%d\n", data);`
- with string(3) functions, particularly
`message = strcat(message, "some string");`
- send with
`sent = send(c, buffer, strlen(buffer), 0);`
- check for errors `sent == -1`

Example

```
const size_t msgsize = 4096;
char msgbuff[msgsize];
char *msg;

strcpy(msgbuf, "Messsage Header\r\n");
strcat(msgbuf, "v1.1\r\n");
msg = strchr(msgbuf, '\0');
msg += sprintf(msg, "Key:%d\r\n", value);

sent = send(con, msgbuff, strlen(msgbuff), 0);
```

Receiving Messages

- declare buffer as `char buffer[BUFSIZE]`
- byte count `size_t bytes`
- receive with
`bytes = recv(c, buffer, BUFSIZE-1, 0);`
- check for errors `bytes == -1`
- add terminating zero for string manipulation
`buffer[bytes]='\0';`

Example

```
const size_t bufsize=4096;
char msgbuf[bufsize];
size_t bytes;

bytes = recv(con, msgbuf, bufsize-1, 0);
if( bytes == -1 ) { /*handle error*/ }
msgbuf[bytes]='\0';
```

Unpacking messages

- Typically messages are line oriented with carriage-return line-feed characters (CRLF) as line terminators
- Some protocols are permissive on what counts as a line ending
- lines are formatted, often key:value pairs.

strtok(3)

The `strtok` pair of functions are ideal for splitting the message into lines and parts.

Splitting a message into lines

Reentrant version of `strtok`

```
/* assume char *message is full message */
char *line;
char *rest;
for(
    line = strtok_r(message, "\r\n", &rest);
    line!=NULL;
    line = strtok_r(NULL, "\r\n", &rest);
) {
    /* handle line in here
}
```

Need reentrant version to remember where we are in `rest`

Beware

As used here, `strtok_r` will skip blank lines. Not helpful if you want to identify the end of an HTTP header.

Split a message into key:value pairs

plain strtok

```
char *key;  
char *value;  
key   = strtok(line, ":");  
value = strtok(NULL, ":");
```

strtok remembers it's place internally, which means it can't be nested.

Creating a server using TCP

- Create server address with 'listening port' `getaddrinfo(3)`
- create a socket `socket(2)`
- `bind(2)` the socket to the address
- tell the socket to `listen(2)`
- use `accept(2)` to listen for connections
- use `send(2)` and `recv(2)` to send and receive messages.

Example

```
int s = socket(AF_INET, SOCK_STREAM, 0);

hints.ai_family = AF_INET;
hints.ai_socktype = SOCK_STREAM;
hints.ai_flags = AI_PASSIVE;

err = getaddrinfo( NULL, "65421", &hints, &results);
err = bind(s, results->ai_addr, results->ai_socklen);
err = listen(s, 1);

int cfd;
struct sockaddr client;
socklen_t size = sizeof(client);
cfd = accept(s, &client, &size);

send and receive on file-descriptor cfd
```


Notes & hints

- *Always* check return values for errors
 - see `errno(3)`
 - see `perror(2)`
- Assume text (for now)
- *Add* zero terminating byte to string in read buffer

how do I know if I have the whole message? The `PROTOCOL` defines the beginning and end of a message.

my read buffer isn't big enough! You'll get the message in parts

- bigger buffer
- mechanism to assemble message from multiple calls to `recv`